

CS220 - Computer System II
Lab 10 – return-to-libc

Due: 11/09/2016, 11:59pm

1 Introduction

In this lab, you will write a program that will spawn a shell, but with a twist. You will not call a function that will actually create the shell, but you will corrupt the return address to *return* to the function that creates a shell. Attacks that invoke functionality by returning to beginning of a function rather than calling a function are called return-to-libc attacks. The libc here is only due to historic reasons, but the attack itself works on any function. `libc` is a library that contains several useful juicy functions worth returning to!

Note: the instructions in this lab will refer to addresses in our address space when we compiled and ran the lab. Your addresses may differ due to the different versions of `libc` or Address Space Layout Randomization (ASLR). These differences do not affect the idea of this lab, but will require your attention to detail.

2 Getting Started

Create a folder **Lab10** with four files: **prelim.c**, **ret2libc.c**, **get_ebp.S** and a Makefile. The Makefile should have three targets:

prelim This executable will demo key ideas discussed in §3.

ret2libc This executable will build your attack using code found in both **ret2libc.c** and **get_ebp.S**.

clean Removes the two above executables and any temporary files. Exactly what you're expecting.

3 Background

Function `int system (char *command)`¹ is a powerful libc function that allows execution of commands from within a C program. Try the following code in **prelim.c**, build it with the Makefile, and you will see the directory listing:

¹Make it a habit to check the man pages whenever you encounter a new library function (i.e. `$ man system`).

```
1 #include <stdlib.h>
   int main() {
3   return system("ls");
   }
```

Similarly, replace "ls" with "/bin/bash", and you will spawn a shell. So, the goal is to invoke `system("/bin/bash");` without actually invoking it.

Strategy: First, we will identify the stack layout when the control enters `system`. Then, we will corrupt the return address of the function and replace it with address of the `system` function. At the same time, we will set up the stack in such a way that when the function returns, it is equivalent to calling `system` with "/bin/bash" as its argument.

4 Implementation

Our first task is to implement a function that will return the value of the `ebp` register as a 32-bit `uint32_t` value (this type is defined in `<stdint.h>`). Add the function prototype `extern uint32_t* get_ebp(void);` to `ret2libc.c`, and implement it in `get_ebp.S`. (Hint: This is a very simple assembly function of only two instructions.)

Our next task is to examine the stack layout when the control enters `system`. Compile `prelim.c` (do not forget the `-m32` flag in your Makefile) and start it in `gdb`. Insert a breakpoint at `system`:

```
1 (gdb) b system
```

When you hit the breakpoint, examine the contents of the stack. Command `x/16x $esp` tells `gdb` to print 16 DWORDS in hex starting from `$esp`.

```
1 (gdb) x/16x $esp
```

Remember `esp` grows towards lower address. So, this will actually give you the first 16 DWORDS on the stack. The first DWORD must be the return address where `system` will return to. Confirm that it is within the `main()` function.

```

1  (gdb) x/16x $esp
   0xffffdb3c: 0x08048431  0x080484d0  0xf7fee560  0x0804845b
3  0xffffdb4c: 0xf7fa4ff4  0x08048450  0x00000000  0xffffdbd8
   0xffffdb5c: 0xf7e57e46  0x00000001  0xffffdc04  0xffffdc0c
5  0xffffdb6c: 0xf7fdc860  0xf7ff47f1  0xffffffff  0xf7ffcff4
   (gdb) disas main
7  Dump of assembler code for function main:
   0x0804841c <+0>: push    %ebp
9  0x0804841d <+1>: mov     %esp,%ebp
   0x0804841f <+3>: and     $0xffffffff0,%esp
11 0x08048422 <+6>: sub     $0x10,%esp
   0x08048425 <+9>: movl    $0x80484d0,(%esp)
13 0x0804842c <+16>: call    0x8048300 <system@plt>
   0x08048431 <+21>: leave
15 0x08048432 <+22>: ret
   End of assembler dump.
17 (gdb)

```

As expected, 0x08048431 is the first entry, and that is the address in main that system returns to. Next entry must be the argument to system. Confirm it.

```

1  (gdb) x/s 0x080484d0
   0x080484d0:  "ls"

```

So, now we know what the stack layout must be. Within ret2libc.c, write a function `void ret2libc()` that will implement your code.

```

1  #include <stdio.h>
2  #include <stdlib.h>
3
4  extern uint32_t* get_ebp();
5
6  void ret2libc(int dummy) {
   /* TASK 1: Get the value of ebp. Make use of the get_ebp function you have
      previously written. */
8  /* TASK 2: Overwrite written address with address of system */
   /* TASK 3: Set up the argument to system as a pointer to string "/bin/bash"
      */
10 }

```

```

12 int main() {
    ret2libc(0);
14     printf("Done!\n");
    return 0;
16 }

```

TASK 1: Notice how we have declared that `get_ebp` returns a pointer to an unsigned 32 bit value. This allows us to treat the return value as a base of an array and access each DWORD on the stack as an array access. For example, suppose the return value of `get_ebp` is `curr_ebp`, then `curr_ebp[0]` is the old `ebp` of the previous frame that was saved using the `push ebp` instruction. Similarly `curr_ebp[1]` is the address where `ret2libc` returns to, and `curr_ebp[2]` is the dummy variable.

TASK 2: Now, go ahead and replace the return address of `ret2libc` with `&system`.

```
curr_ebp[1] = &system; /* Compiler will complain. Fix it. */
```

TASK 3: We know that when control enters `system`, the first DWORD will be the address that `system` returns to, but the second DWORD is where the argument is stored. So, create a string that contains the argument, and set it to `curr_ebp[3]`:

```

1 char *str = "/bin/bash";
  curr_ebp[3] = str;

```

Compile and run, and you should get a bash prompt. The `ret2libc()` function actually “returned” to `system`. The `system` function promptly fed off of the arguments on the stack, and executed a shell. Go ahead and exit the shell.

```

$ exit
2 exit
Segmentation Fault

```

The program crashed because the `system` function tried to return to whatever was present in `curr_ebp[2]`, which was not to where `ret2libc()` was originally to return. In fact, `ret2libc` must return to `main`. In order to fix it, we simply copy the original return address

of `ret2libc` function to `curr_ebp[2]`. This should happen at the beginning before we replace it with address of `system`.

```
1 curr_ebp[2] = curr_ebp[1];  
...
```

Compile and run the program. You should get a shell, and upon exiting from the shell, you should see “DONE!” printed from `main`.

5 Adding a little spice!

Now, based on what you have learned so far, implement `void ret2libc_generic(char *cmd);`. Your goal is to supply `cmd` as argument to `system`, as opposed to a fixed string “/bin/bash” in the previous case. The challenge here is that `cmd` resides at `curr_ebp[2]`. So, be careful when you overwrite it! Move it first to the location where argument to `system` goes.

6 Submitting the result

Once you have implemented and tested `ret2libc_generic()`, you are ready to submit your code. Remove binaries and intermediate files from `Lab10/`. Create a `tar.gz` of `Lab10/` folder with only the files listed in §2.

```
$ tar -cvzf lab12_submission.tar.gz ./Lab10
```

Submit `lab10_submission.tar.gz` to MyCourses.